

Machine learning notes: week 3

Lars Yencken

Logistic regression

Classification

- When the outcome we're predicting is boolean (or, an enum)
- Examples
 - Is this email spam?
 - Is this tumour malignant?
 - Has the deadly neurotoxin worked?
- Ghetto approach: apply linear regression and use some threshold, but...
 - Outlying examples can really screw up the model
 - Outputs can be way outside the range $[0, 1]$
- Better approach: use a different family of models which are S shaped curves instead of straight lines

Hypothesis representation

- Let $g(z) = \frac{1}{1+e^{-z}}$ where $z = \theta^T x$
- g is called the *{sigmoid function} or the {logistic function}*
- In practice, this function squishes any z value into something between 0 and 1
 - Large negative values of z become close to zero
 - Large positive values of z become close to one
- Now that $h_\theta(x)$ is between 0 and 1, call it the "probability that $y = 1$ "
 - Formally $h_\theta(x) = \Pr(y = 1|x; \theta)$

Decision boundary

- We want to decide if y is 0 or 1, but $h_\theta(x)$ is a floating point number between 0 and 1
- If we pick 0.5 as a threshold, then it means we'll pick
 - $y = 1$ if $\theta^T x > 0$
 - $y = 0$ if $\theta^T x < 0$
- A threshold like this cuts the space for x into two parts, where we always predict $y = 1$ on one side and $y = 0$ on the other
- It's called a *decision boundary*
- Decision boundaries can be non-linear, if we have non-linear features (e.g. polynomial features)

Cost function

- Given our training set, how can we pick θ ?
- Use a cost function to determine which is best (i.e. lowest cost), gradient descent to find that θ
- The naive translation of $J(\theta)$ is non-convex (i.e. full of local minima)
- Instead, let $cost(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y = 1 \\ -\log(1 - h_\theta(x)) & \text{if } y = 0 \end{cases}$
- This adds a generous penalty for being wrong in any individual case, and will make $J(\theta)$ convex
- Once again, you don't need to be able to come up with this math, it's enough to understand roughly what it's doing

Gradient descent

- We can add up costs for each case in the training set to get $J(\theta)$:
 - $J(\theta) = \frac{1}{m} \sum_{i=1}^m cost(h_\theta(x^{(i)}, y^{(i)}))$
- We can simplify the cost component to one line:
 - $J(\theta) = \frac{1}{m} [\sum_{i=1}^m y^{(i)} \log(h_\theta(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)}))]$
- For gradient descent, rule looks very similar to linear regression:
 - $\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$
 - Notice that $h_\theta(x)$ is the only part that's changed
- The same rules as normal apply:
 - Monitor $J(\theta)$ and check that it's decreasing with more iterations
 - Try various values for α and tune appropriately

Advanced optimisation

- Gradient descent calculates $J(\theta)$ and $\frac{\delta}{\delta \theta_j} J(\theta)$ in every loop
- A bunch of advanced algorithms exist which can:
 - Avoid the need to pick α manually
 - Can converge much faster than gradient descent
- They're quite complex to implement – don't roll your own!
- These become quite important for scaling to larger datasets

Multiclass problems

- `multiclass`: predicting an enum instead of a boolean
- Examples
 - Email tagging: work, friends, family or hobby
 - Medical diagnosis: not ill, cold, flu
- One vs all
 - Train a model to detect each class, ending up with n models
 - When predicting, run it against all the models and pick the most confident